

Reachable-donor future transplantation

Do generated futures causally and directionally control robot behavior?

— recipient continuation A — donor continuation B - - - targeted future channel / replacement — simulator execution □□□ token sequence ■ learned model

1 Pair native continuations

🔒 same restored simulator state S

A · recipient



B · donor

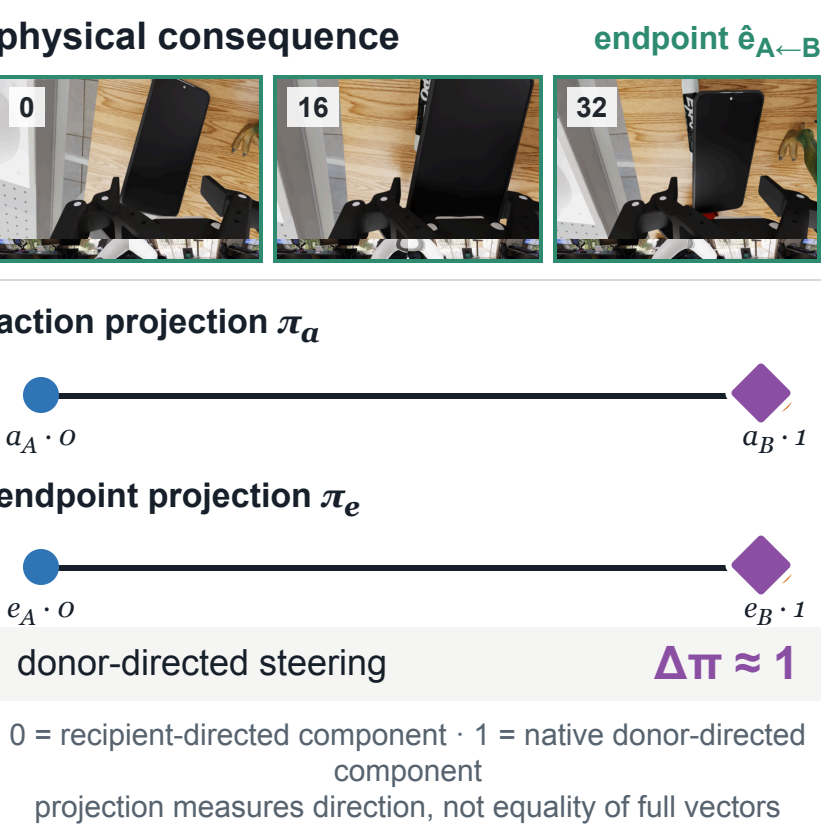


Both branches are policy-generated and physically reachable. Their native actions define a direction; they are not candidate actions shown to the model

2 Transplant only the future



3 Execute and score



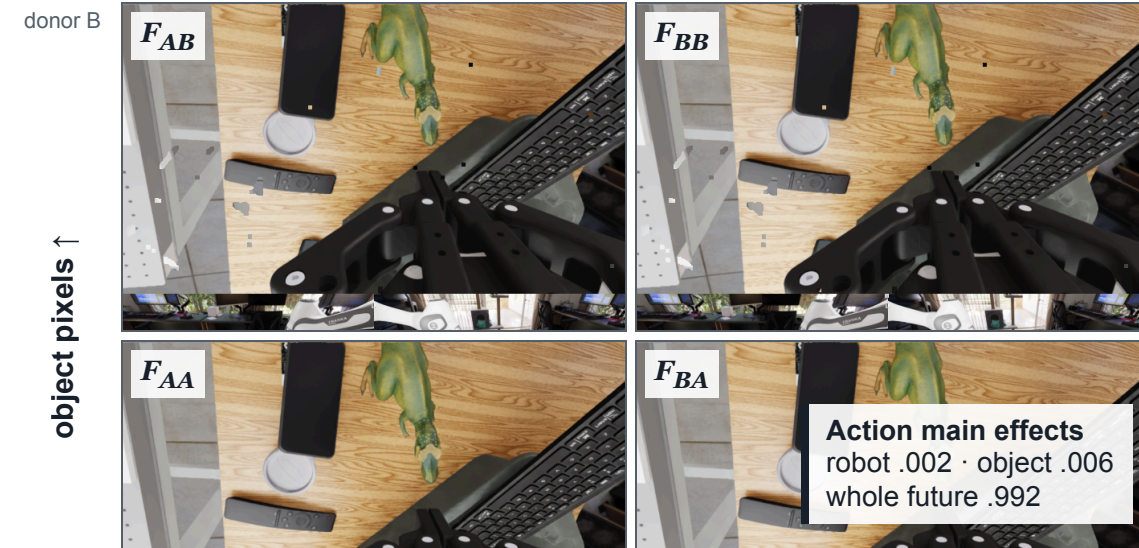
Q1 Does coherent future content steer behavior?

Cosmos 3 · 22 restored states · six RoboLab tasks



Q2 Are isolated robot or object pixels sufficient?

Actual 2x2 hybrid futures · object-priority disjoint masks



Q3 Which internal route carries donor influence?

Cosmos 3 · token-count-preserving attention intervention

