

---

# How Do World Action Models Use Their Futures?

## Evidence from Mechanistic Interventions

---

Nicole Ma

### Abstract

1 Do world action models actually use the futures they generate? Jointly producing  
2 future observations and robot actions does not answer this question: the action  
3 may bypass the predicted future. We test future use with reachable-donor trans-  
4 plantation. Two native policy runs begin from the same restored simulator state  
5 and produce different futures, actions, and endpoints. We call the run being recom-  
6 puted the recipient and the run supplying the alternative future the donor. Holding  
7 the recipient’s observation, instruction, sampling noise, and denoising schedule  
8 fixed, we ask not merely whether its action changes, but whether it moves specifi-  
9 cally toward the donor action. In Cosmos 3, a frozen study across 22 states and six  
10 RoboLab tasks finds predicted-future action steering of 0.992 and executed-future  
11 endpoint steering of 1.003 on a scale where 0 is the recipient and 1 is the donor.  
12 Replacing donor-run future-token keys and values with self-future counterparts  
13 removes 0.877 of the predicted-action shift, with positive action and endpoint re-  
14 ductions in all 22 states. Isolated robot- and object-pixel interventions, however,  
15 each have effects below 0.011. Cosmos Policy shows a related early-state effect  
16 across all ten LIBERO-Long tasks, carried mainly by visual content. These results  
17 show that both direct policies use coherent futures to condition action, but do not  
18 establish candidate search or planning over predicted outcomes.

## 19 1 Introduction

20 Many world action models (WAMs) generate future observations alongside robot actions [8, 14]. It  
21 is tempting to read this architecture literally: the model predicts what may happen, then acts on that  
22 prediction. Joint generation does not establish such a computation. The action could be generated  
23 through a direct observation-to-action route while a plausible video is produced alongside it.

24 This leaves at least four hypotheses about the future’s role: an epiphenomenal output with no func-  
25 tional role; a generic workspace whose presence matters but whose specific content does not; a  
26 content-specific action interface; or a representation of consequences used to compare actions. The  
27 third role includes an inverse-dynamics-like computation in which a prospective robot pose specifies  
28 the action needed to produce it [1, 7]. The fourth more closely resembles explicit visual planning:  
29 generate alternatives, predict their outcomes, score them with a cost or learned heuristic, and se-  
30 lect among them [2, 3]. We study Cosmos Policy’s direct-policy mode and the diffusion-generator  
31 component of Cosmos3-Nano-Policy-DROID. Neither evaluated inference procedure generates and  
32 scores multiple candidate actions before acting; Cosmos Policy’s separate best-of- $N$  planning mode  
33 is outside our study.

34 RIFT provides the closest causal evidence that future representations matter [22]. Masking the fu-  
35 ture read or editing values stored at future positions changes closed-loop trajectories and task success  
36 across several WAMs. Such destructive tests establish that the future channel is necessary or behav-  
37 iorally sensitive. They do not establish that its semantic content controls a particular alternative: a  
38 damaged representation can move an action in any direction.




Figure 1: Reachable-donor future transplantation. (a) Two native continuations start from the same restored state. (b) The recipient is recomputed with only the donor future substituted. (c) The donor action is withheld and used afterward to score signed movement from recipient (0) toward donor (1). (d) Replacing donor-future K/V with self-future K/V tests whether that interface carries the directional shift.

Our test is directional. From the same restored state, two native policy continuations define a recipient and a donor. Both are coherent alternatives generated by the unmodified policy, and their one-chunk endpoints are reachable from that state. We insert the donor future into the recipient computation while retaining the recipient’s current inputs and random draws. The model never receives the donor action; we use it only afterward to define the direction in which we score the recomputed action and its executed simulator endpoint (Figure 1). A generic perturbation need not move either quantity toward that particular donor. Directional movement therefore tests content-specific control, rather than dependence on the future channel alone.

We ask three questions: Do coherent donor futures steer behavior? Which parts of the future recreate that effect? Which future-to-action interfaces carry it? Cosmos 3 provides the strongest within-model chain: steering generalizes across 22 sampled states, isolated robot or object pixels do not reproduce it, and future-token K/V replacement suppresses action and endpoint steering in every state. Cosmos Policy supplies an earlier-model comparison whose effects are concentrated at early states and are most consistent with camera-visible robot motion. The evidence supports future-dependent action generation, but not candidate search, outcome comparison, or task-level planning.

## 2 Related Work

**Prediction and control.** Some visual robot policies generate future frames and infer actions from them; others condition action inference on predicted video representations [1, 7]. Joint video–action models instead learn future observations and controls in one model [6, 8, 10, 14, 20]. Their training objective does not determine whether a particular generated future is used at inference. UVA, for example, can omit video generation at test time [10]. Explicit visual planners answer a different question by sampling alternatives, predicting their outcomes, and comparing them with a goal-dependent score [2, 3].

Recent WAM studies separate predictive training from the use of predictions during action generation. Fast-WAM removes explicit test-time rollout while retaining competitive performance in its evaluations [21]; Faster-WAM retains a lower-cost future-conditioning route and reports improved

robustness [24]. AGRA finds that visually plausible foresight need not support reliable action inference in its studied architecture [16]. RIFT directly establishes dependence on future-position values and their organization [22]. Our contribution is complementary but distinct: RIFT asks whether behavior depends on the future interface, whereas reachable-donor transplantation asks whether a coherent alternative sent through that interface predicts the direction of behavior.

**Mechanistic interventions.** Activation replacement tests whether an internal representation affects an output rather than merely being decodable from it [12, 18]. Path patching extends this logic to specified computational routes [5]. Conclusions can depend on the counterfactual pair and metric [23], while multi-component effects need not add when patched and unpatched routes interact [17]. We therefore treat K/V replacement as evidence for an active pathway, not a complete circuit decomposition.

## 3 Methods

### 3.1 Future transplantation and directional projection

Let saved simulator state  $S$  admit two unmodified branches  $A$  and  $B$ , with future representations  $F_A, F_B$ , action chunks  $\mathbf{a}_A, \mathbf{a}_B$ , and endpoints  $\mathbf{e}_A, \mathbf{e}_B$  obtained by executing one action chunk. We call  $A$  the recipient—the branch whose action is recomputed—and  $B$  the donor—the branch supplying the alternative future. To recompute  $A$ , we restore the same state and observation, retain its instruction, action noise, future-noise path, solver, and denoising schedule, and replace only  $F_A$  with  $F_B$ . We repeat the intervention in the reciprocal direction, transplanting  $F_A$  into  $B$ .

For an action or endpoint  $\mathbf{x} \in \{\mathbf{a}, \mathbf{e}\}$ , normalized donor projection is

$$\pi_x(\hat{\mathbf{x}}; A, B) = \frac{(\hat{\mathbf{x}} - \mathbf{x}_A)^\top (\mathbf{x}_B - \mathbf{x}_A)}{\|\mathbf{x}_B - \mathbf{x}_A\|_2^2}. \quad (1)$$

The recipient is 0, the donor is 1, and 0.5 is halfway between them along the donor direction; negative values move away and values above 1 overshoot. Because the dot product retains only this signed direction, a large change perpendicular to the donor receives little projection. The primary steering contrast subtracts the corresponding self-future recomputation:

$$\Delta_x = \pi_x(\hat{\mathbf{x}}_{\text{donor}}; A, B) - \pi_x(\hat{\mathbf{x}}_{\text{self}}; A, B). \quad (2)$$

The self condition reruns the same intervention machinery with the recipient’s own future, so subtraction removes any displacement due merely to recomputation. Pairs are selected using native action and endpoint separation, never intervention outcomes. Every endpoint arm is executed after restoring the exact simulator state.

### 3.2 Models, populations, and controls

**Cosmos Policy.** We evaluate the public `Cosmos-Policy-LIBERO-Predict2-2B` checkpoint in direct-policy mode on LIBERO-Long, also called LIBERO-10 [8, 11]. The frozen grid contains 30 states: ten tasks crossed with the first query, a state after three open-loop chunks, and a state after six chunks. Eight native runs determine one reciprocal donor pair per state; three future-noise draws are repeated within direction. Endpoints concatenate goal-relevant simulator coordinates with the official nine-dimensional proprioceptive endpoint.

**Cosmos 3.** We evaluate the 16B `Cosmos3-Nano-Policy-DROID` checkpoint on six eligible RoboLab tasks [14, 15, 19]. Its diffusion generator has 36 two-way full-attention layers, four denoising calls, 33 raw video frames mapped to nine VAE latent frames, and 32 absolute joint-position actions. The transplant targets unconditioned future-video coordinates at the final guided-sampler boundary. A frozen screen selected 24 state clusters, capped at four per task, using only native success, action and endpoint separation, and replay validity. Twenty-two completed. Two source branches ended before producing full 33-frame videos and were left missing rather than replaced. Reciprocal directions and predicted versus executed future sources are repeated measurements within state.

Controls include exact self-future recomputation, Gaussian targets matched to the donor in norm and recipient distance, a prespecified natural-future control, within-task shuffled futures, and bitwise no-op checks. Gaussian controls test whether an equally large unstructured displacement is sufficient; the natural and shuffled controls compare the paired donor against a prespecified non-donor continuation and mismatched futures. For Cosmos 3, identity layout capture, sampler wrapping, cache record/replay, and self patches reproduce actions bitwise; sentinels confirm that the transplant does not directly write to action coordinates. Executed-future conditions encode videos produced by physically executing reachable donor actions and replay the same environment prefix before intervention. They test whether steering survives a physically realizable continuation rather than appearing only for model-predicted imagery.

### 3.3 Robot and object interventions

We isolate content in complementary ways. In Cosmos Policy, first-query interventions replace the wrist-camera future, primary-camera future, or future proprioception separately. A second study crosses binary robot and object states in re-rendered videos at 20 held-out states. A decoder checks whether the intended cell remains visible. A natural-pair screen separately searches policy-generated continuations for robot-divergent/object-matched and object-divergent/robot-matched futures; only four states meet its criteria, below the planned minimum of ten.

For Cosmos 3, a separately frozen factor study selected 12 population states round-robin by native object-position separation. Ten completed across all six tasks. Starting from the recipient video, an object-priority disjoint mask copies donor object pixels, donor robot pixels outside the object mask, both, or neither. This  $2 \times 2$  design tests isolated pixel factors; it does not guarantee an in-distribution video or isolate a latent semantic variable.

### 3.4 K/V interventions

In attention, keys help determine what an action query reads, while values carry the information returned by that read. In Cosmos Policy, action queries can attend to future keys. We continuously gate future keys at prespecified block 27, with exact all-key recomputation, equal-count random-key and current-key controls, and a block-0 control. This removal changes attention normalization, so it tests whether a late future-to-action edge is active, not the information carried by a token-count-preserving route.

Cosmos 3 permits the stronger replacement test. We record future-video keys and values from a self-future run, then substitute them for donor-run future-token K/V at the frozen direct interface. Token count, token order, current-video K/V, text K/V, action K/V, observation, instruction, and noise remain fixed. Layer and pathway choices were frozen on excluded calibration tasks before the 22-state population intervention. We report *K/V replacement loss*: the reduction in donor projection after substituting self-future K/V. This is an operational intervention contrast, not a population natural indirect effect.

### 3.5 Statistical analysis

The independent unit is a restored simulator state, not an individual noise draw or transplant direction. We first average reciprocal directions, donors, noise draws, and process repeats within state, then form 95% percentile intervals from 10,000 bootstrap resamples of states [4]. For the common-scale condition estimates in Table 1, we reconstruct each control’s projection within unit from the registered pairwise contrasts before bootstrapping; these are not differences of aggregate means. Task-balanced sensitivity analyses average states within task before resampling tasks. We prespecified 0.10 projection units—10% of the native recipient-to-donor displacement—as the smallest effect of interest for the Cosmos Policy confirmatory family and the Cosmos 3 population and factor families [9]. An interval contained within  $[-0.10, 0.10]$  supports practical equivalence to zero at that scale. Prespecified primary sign-test  $p$ -values use Holm correction.

Table 1: Method and controls on one scale. Each row reports donor-directed projection after inserting the named future: 0 is the self-future result and 1 is the paired donor behavior. Within each model–source–outcome block, bold marks the largest estimate, always the paired donor (ours). The rightmost column gives its paired advantage over the strongest available control; all five intervals exclude zero. Control projections are reconstructed within independent unit before bootstrapping.

| Model         | Source    | Outcome  | Future inserted            | Projection [95% CI]         | Best-control advantage [95% CI] |
|---------------|-----------|----------|----------------------------|-----------------------------|---------------------------------|
| Cosmos 3      | Predicted | Action   | Self                       | 0.000 (reference)           | —                               |
| Cosmos 3      | Predicted | Action   | Natural control            | 0.559 [0.465, 0.651]        | —                               |
| Cosmos 3      | Predicted | Action   | <b>Paired donor (ours)</b> | <b>0.992 [0.972, 1.010]</b> | <b>0.433 [0.332, 0.538]</b>     |
| Cosmos 3      | Executed  | Action   | Self                       | 0.000 (reference)           | —                               |
| Cosmos 3      | Executed  | Action   | Matched Gaussian           | 0.041 [0.003, 0.076]        | —                               |
| Cosmos 3      | Executed  | Action   | Natural control            | 0.591 [0.491, 0.687]        | —                               |
| Cosmos 3      | Executed  | Action   | <b>Paired donor (ours)</b> | <b>0.983 [0.890, 1.064]</b> | <b>0.391 [0.246, 0.533]</b>     |
| Cosmos 3      | Executed  | Endpoint | Self                       | 0.000 (reference)           | —                               |
| Cosmos 3      | Executed  | Endpoint | <b>Paired donor (ours)</b> | <b>1.003 [0.950, 1.050]</b> | <b>1.003 [0.950, 1.050]</b>     |
| Cosmos Policy | Predicted | Action   | Self                       | 0.000 (reference)           | —                               |
| Cosmos Policy | Predicted | Action   | Matched Gaussian           | 0.001 [−0.004, 0.006]       | —                               |
| Cosmos Policy | Predicted | Action   | Natural control            | 0.395 [0.269, 0.553]        | —                               |
| Cosmos Policy | Predicted | Action   | Shuffled control           | 0.244 [0.152, 0.347]        | —                               |
| Cosmos Policy | Predicted | Action   | <b>Paired donor (ours)</b> | <b>0.499 [0.336, 0.679]</b> | <b>0.103 [0.048, 0.161]</b>     |
| Cosmos Policy | Predicted | Endpoint | Self                       | 0.000 (reference)           | —                               |
| Cosmos Policy | Predicted | Endpoint | Matched Gaussian           | 0.003 [−0.003, 0.008]       | —                               |
| Cosmos Policy | Predicted | Endpoint | Natural control            | 0.442 [0.311, 0.598]        | —                               |
| Cosmos Policy | Predicted | Endpoint | Shuffled control           | 0.275 [0.185, 0.375]        | —                               |
| Cosmos Policy | Predicted | Endpoint | <b>Paired donor (ours)</b> | <b>0.552 [0.387, 0.728]</b> | <b>0.110 [0.054, 0.166]</b>     |

## 4 Results

### 4.1 Do coherent futures steer behavior?

The method succeeds if three things happen: inserting a donor future moves behavior toward that donor, the shift survives simulator execution, and the paired donor outperforms non-donor controls. Table 1 puts the method and every applicable control on one scale, organized by model, future source, and outcome. Figure 2 shows the same comparison visually. On this scale, 0 is the self-future result and 1 is the paired donor behavior.

**Cosmos 3.** Across 22 completed states and all six tasks, a predicted paired-donor future produces action steering of 0.992 ([0.972, 1.010]; 22/22 states positive), compared with 0.559 for the natural-future control (Table 1). Executed paired-donor videos produce action steering of 0.983, compared with 0.591 for the natural control and 0.041 for the geometry-matched Gaussian control. Executing the paired-donor-conditioned action yields endpoint steering of 1.003 ([0.950, 1.050]; 22/22). Thus, the effect is not explained by recomputation, an equally large unstructured displacement, or merely supplying a different valid continuation. Endpoint projection shows that the donor-directed component survives simulator execution; it does not imply equality of the full endpoint vectors. Task-balanced analyses preserve these conclusions.

**Cosmos Policy.** At the first query, the paired donor produces action steering of 0.499 and endpoint steering of 0.552, positive in all ten tasks (Table 1). The natural control reaches 0.395 and 0.442, respectively, while the shuffled control reaches 0.244 and 0.275 and the matched Gaussian remains near zero. The natural future is the hardest control: the paired donor exceeds it by 0.103 ([0.048, 0.161]; 8/10) for actions and 0.110 ([0.054, 0.166]; 8/10) for endpoints. This rules out the weak explanation that any different tokenizer-consistent future causes the same movement. Effects are much smaller at the separately sampled later states: after three chunks, action steering is 0.031 ([0.006, 0.066]) and endpoint steering is 0.004 ([−0.043, 0.038]); after six, they are 0.027 ([−0.010, 0.092]) and 0.025 ([0.005, 0.044]). Because each time point is a different physical state, this is state dependence associated with time, not a causal estimate of temporal decay.




Figure 2: Method and controls on the same donor-directed scale, separated by model. Blue squares are self-future baselines; gray marks are Gaussian, natural, and shuffled controls; orange diamonds are paired donor futures (ours). Action and endpoint are separate x-axis groups, not marker meanings. Points show means and whiskers show 95% state-bootstrap intervals. In every group, the paired donor has the largest projection. The dotted orange line marks the paired donor behavior at 1.

182 No screened branch pair in either policy retained a stable success/failure label across continuation  
 183 seeds. We therefore interpret these results as donor-directed changes in action and one-chunk end-  
 184 point, not changes in closed-loop task completion.

#### 185 4.2 What future content causes steering?

186 The content that drives steering differs across checkpoints. Cosmos Policy responds mainly to  
 187 camera-visible robot motion at early states, whereas neither isolated pixel factor recreates Cosmos  
 188 3’s whole-future effect.

189 **Cosmos Policy.** Wrist-future action steering is 0.435 ([0.284, 0.605]), compared with 0.101  
 190 ([0.057, 0.144]) for the primary camera and 0.001 ([−0.004, 0.008]) for future proprioception (Fig-  
 191 ure 3a). In the 20-state  $2 \times 2$  study, decoded futures identify the intended cell with 92.5% primary-  
 192 camera and 98.75% wrist-camera accuracy. The intended object manipulation is therefore visible  
 193 in the rendered futures, yet the object-state main effect is 0.00015 ([−0.00012, 0.00049]) for ac-  
 194 tion and 0.00018 ([−0.00040, 0.00102]) for the executed goal-state endpoint. The available analysis  
 195 does not report the robot-state main effect. Four exploratory natural pairs provide a complementary  
 196 comparison: robot-motion futures yield action steering of 0.291 ([0.060, 0.521]) and robot-endpoint  
 197 steering of 0.315 ([0.059, 0.571]), whereas the object-only effects are 0.021 for both action and goal  
 198 endpoint, with intervals crossing zero. At these sampled early states, the pattern is most consistent  
 199 with an inverse-dynamics-like role for camera-visible robot motion. It does not show that Cosmos  
 200 Policy generally ignores object consequences.

201 **Cosmos 3.** Across ten factor-study states and all six tasks, every isolated-factor interval lies in-  
 202 side the prespecified [−0.10, 0.10] equivalence region (Table 2; Figure 3b). Robot-pixel effects are  
 203 0.0022 for actions and 0.0102 for endpoints; object-pixel effects are 0.0061 and −0.0001. Neither  
 204 factor alone recreates a meaningful fraction of the roughly one-unit whole-future effect. The use-  
 205 ful signal may be distributed across the image, require robot–object coherence, depend on camera  
 206 or temporal structure, or interact nonlinearly with masking and tokenization. Hybrid videos may  
 207 also be off distribution. The experiment establishes insufficiency of each isolated pixel factor, not  
 208 irrelevance of robot or object information and not joint necessity.



Table 2: Cosmos 3 isolated content factors across ten states. Each row is a main effect in the disjoint  $2 \times 2$  pixel transplant. Every 95% state-bootstrap CI is contained within the prespecified equivalence region.

| Factor        | Outcome  | Effect [95% CI]           | Equivalent to zero |
|---------------|----------|---------------------------|--------------------|
| Robot pixels  | Action   | 0.0022 [−0.0016, 0.0060]  | Yes                |
| Robot pixels  | Endpoint | 0.0102 [−0.0067, 0.0275]  | Yes                |
| Object pixels | Action   | 0.0061 [0.0009, 0.0119]   | Yes                |
| Object pixels | Endpoint | −0.0001 [−0.0254, 0.0239] | Yes                |

Table 3: Cosmos 3 future-token pathway across 22 states. Each unpatched donor-minus-self effect is shown beside the reduction caused by replacing donor-future K/V with self-future K/V. “Removed” is the ratio of point estimates, not an independently estimated mediation proportion.

| Outcome           | Unpatched steering   | K/V replacement loss | Point estimate removed |
|-------------------|----------------------|----------------------|------------------------|
| Predicted action  | 0.992 [0.972, 1.010] | 0.877 [0.842, 0.910] | 88%                    |
| Executed action   | 0.983 [0.890, 1.064] | 0.814 [0.723, 0.900] | 83%                    |
| Executed endpoint | 1.003 [0.950, 1.050] | 0.851 [0.766, 0.922] | 85%                    |

### 4.3 Which internal pathway carries the effect?

Future-token attention provides an active route from generated futures to action. Population-level K/V replacement in Cosmos 3 gives the stronger of the two pathway tests.

**Cosmos Policy.** Block-27 future-key removal produces a monotonic action-disruption curve: normalized  $L_2$  change is 0.022, 0.035, 0.054, and 0.068 at gates 0.25, 0.50, 0.75, and 1.00. All-key recomputation is bitwise exact. Full removal is 0.016 more disruptive than equal-count random-key removal and 0.048 more disruptive than block-0 future-key removal; equal-count current-key removal is 0.024 larger. Executed endpoints support the same bounded conclusion, although state counts and intervals are not available for this analysis. The late future-to-action edge is active, but it is neither the only nor the largest route through the block.

**Cosmos 3.** The interface was selected on excluded calibration tasks and frozen before population testing. Replacing donor-run future-token K/V with self-future K/V removes 0.877 of the 0.992 predicted-action shift, 0.814 of the 0.983 executed-action shift, and 0.851 of the 1.003 endpoint shift (Table 3; Figure 3c). These losses correspond to 88%, 83%, and 85% of the respective unpatched point estimates and are positive in 22/22 states. Task-balanced point estimates are 0.877, 0.812, and 0.853; all prespecified primary sign tests remain significant after Holm correction. Because token count and attention positions remain fixed, donor-specific information travels through this future-token K/V interface. The experiment does not identify every head or feature, exclude interacting unpatched routes, or establish a complete circuit.

## 5 Discussion and Limitations

Across both WAM generations, transplanting a coherent future predicts which alternative action the model will approach. The associated simulator endpoint moves in the same direction. RIFT established that behavior depends on an organized future interface; our result adds content-specific sufficiency by showing that a policy-native alternative predicts the signed direction of the change. The two checkpoints, however, do not appear to rely on the same content.

Cosmos Policy’s early-state pattern resembles an inverse-dynamics interface conditioned on visible robot motion. That account does not transfer cleanly to Cosmos 3: strong whole-future steering coexists with negligible robot-only and object-only pixel effects. Usable content may instead depend on architecture, surrounding context, or distributed feature interactions. The pathway experiments locate where some of this influence travels. Cosmos Policy identifies an active late attention edge, while Cosmos 3 shows across 22 states that future-token K/V carries most of the donor-directed shift.




Figure 3: What carries the effect? Orange denotes donor-future steering, gray denotes isolated pixel-factor effects, and blue denotes the fraction removed by restoring self-future K/V. (a) In Cosmos Policy, the wrist-video future reproduces most of the whole-future action effect; future proprioception reproduces none. (b) In Cosmos 3, neither robot-only nor object-only pixels reproduce the whole-future effect; the shaded band is the prespecified equivalence region. (c) Replacing the donor-future K/V pathway with self-future K/V removes 83–88% of donor steering.

This evidence is narrower than planning. A planning claim would require evidence that the policy generates alternatives, predicts their outcomes, compares them using a goal, cost, or value, and selects accordingly. We do not test that sequence, and no screened pair retained a stable success/failure label across continuation seeds. Here, “future” and “imagined” describe generated representations, not deliberation. Our claim is that these direct policies condition action on coherent futures, not that WAMs generally plan over task consequences.

**Limitations.** All interventions are in simulation and cover two checkpoints, six RoboLab tasks, and LIBERO-Long. Cosmos Policy’s strongest result uses one first-query state per task, while later estimates use different physical states and therefore cannot isolate elapsed time. Its natural content analysis contains only four states, and rendered hybrids may be outside the policy’s native future distribution. Cosmos 3 completed 22 of 24 frozen population states and 10 of 12 factor states; the two population omissions resulted from short source videos and were not replaced. Pixel masks may miss distributed semantic features or induce nonlinear changes, so negligible isolated effects do not establish joint necessity. Key removal changes attention normalization, and K/V replacement can interact with unpatched routes. We do not demonstrate task-success mediation or recover a full circuit.

**Broader impact.** Mechanistic tests can help distinguish plausible-looking predictions from representations that actually control behavior. The same interfaces could also redirect robot actions or create false confidence if validated on narrow state distributions. Use beyond simulation would require broader coverage, physical safety constraints, and monitoring outside the intervention metric.

## 6 Conclusion

The central result is directional: we do not merely damage a model’s future and observe that behavior changes; we supply a different coherent future and find that behavior follows that particular alternative. Cosmos 3 shows this across 22 sampled states, with future-token K/V carrying most of the action and endpoint shift. Cosmos Policy shows a related, more state-dependent effect dominated by visual robot-motion information, while neither isolated robot nor object pixels reproduce Cosmos 3’s whole-future effect. These models use generated futures to condition action. Whether they compare alternative predicted outcomes in order to plan remains unanswered.

## References

- [1] Yilun Du, Sherry Yang, Bo Dai, Hanjun Dai, Ofir Nachum, Joshua B. Tenenbaum, Dale Schuurmans, and Pieter Abbeel. Learning universal policies via text-guided video generation. In *Advances in Neural*



- 272 *Information Processing Systems*, volume 36, 2023. URL [https://proceedings.neurips.cc/  
273 paper\\_files/paper/2023/hash/1d5b9233ad716a43be5c0d3023cb82d0-Abstract.  
274 html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/1d5b9233ad716a43be5c0d3023cb82d0-Abstract.html).
- 275 [2] Yilun Du, Sherry Yang, Pete Florence, Fei Xia, Ayzaan Wahid, Brian Ichter, Pierre Sermanet, Tianhe  
276 Yu, Pieter Abbeel, Joshua B. Tenenbaum, Leslie Kaelbling, Andy Zeng, and Jonathan Tompson. Video  
277 language planning. In *The Twelfth International Conference on Learning Representations*, 2024. URL  
278 <https://openreview.net/forum?id=9pKtcJcMP3>.
- 279 [3] Frederik Ebert, Chelsea Finn, Sudeep Dasari, Annie Xie, Alex Lee, and Sergey Levine. Visual  
280 foresight: Model-based deep reinforcement learning for vision-based robotic control. *arXiv preprint*  
281 *arXiv:1812.00568*, 2018. URL <https://arxiv.org/abs/1812.00568>.
- 282 [4] C. A. Field and A. H. Welsh. Bootstrapping clustered data. *Journal of the Royal Statistical Society: Series*  
283 *B*, 69(3):369–390, 2007. doi: 10.1111/j.1467-9868.2007.00593.x.
- 284 [5] Nicholas Goldowsky-Dill, Chris MacLeod, Lucas Sato, and Aryaman Arora. Localizing model behavior  
285 with path patching, 2023. URL <https://arxiv.org/abs/2304.05969>. arXiv preprint, version  
286 2.
- 287 [6] Yanjiang Guo, Yucheng Hu, Jianke Zhang, Yen-Jen Wang, Xiaoyu Chen, Chaochao Lu, and Jianyu Chen.  
288 Prediction with action: Visual policy learning via joint denoising process. In *Advances in Neural In-*  
289 *formation Processing Systems*, volume 37, pages 112386–112410. Curran Associates, Inc., 2024. doi:  
290 10.52202/079017-3570. URL [https://proceedings.neurips.cc/paper\\_files/paper/  
291 2024/hash/cbe25fa0e7c7084049276888a09acc8d-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/cbe25fa0e7c7084049276888a09acc8d-Abstract-Conference.html).
- 292 [7] Yucheng Hu, Yanjiang Guo, Pengchao Wang, Xiaoyu Chen, Yen-Jen Wang, Jianke Zhang, Koushil  
293 Sreenath, Chaochao Lu, and Jianyu Chen. Video prediction policy: A generalist robot policy with predic-  
294 tive visual representations. In *Proceedings of the 42nd International Conference on Machine Learning*,  
295 volume 267 of *Proceedings of Machine Learning Research*, pages 24328–24346. PMLR, 2025. URL  
296 <https://proceedings.mlr.press/v267/hu25g.html>.
- 297 [8] Moo Jin Kim, Yihuai Gao, Tsung-Yi Lin, Yen-Chen Lin, Yunhao Ge, Grace Lam, Percy Liang, Shuran  
298 Song, Ming-Yu Liu, Chelsea Finn, and Jinwei Gu. Cosmos Policy: Fine-tuning video models for visuo-  
299 motor control and planning. In *The Fourteenth International Conference on Learning Representations*,  
300 2026. URL <https://openreview.net/forum?id=wPEIStHxYH>.
- 301 [9] Daniël Lakens. Equivalence tests: A practical primer for t tests, correlations, and meta-analyses. *Social*  
302 *Psychological and Personality Science*, 8(4):355–362, 2017. doi: 10.1177/1948550617697177.
- 303 [10] Shuang Li, Yihuai Gao, Dorsa Sadigh, and Shuran Song. Unified video action model, 2025. URL  
304 <https://arxiv.org/abs/2503.00200>. arXiv preprint, version 3.
- 305 [11] Bo Liu, Yifeng Zhu, Chongkai Gao, Yihao Feng, Qiang Liu, Yuke Zhu, and Peter Stone. LIBERO:  
306 Benchmarking knowledge transfer for lifelong robot learning. In *Advances in Neural Informa-*  
307 *tion Processing Systems*, volume 36, pages 44776–44791. Curran Associates, Inc., 2023. doi:  
308 10.52202/075280-1939. URL [https://proceedings.neurips.cc/paper\\_files/paper/  
309 2023/hash/8c3c666820ea055a77726d66fc7d447f-Abstract-Datasets\\_and\\_  
310 Benchmarks.html](https://proceedings.neurips.cc/paper_files/paper/2023/hash/8c3c666820ea055a77726d66fc7d447f-Abstract-Datasets_and_Benchmarks.html).
- 311 [12] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and edit-  
312 ing factual associations in GPT. In *Advances in Neural Information Processing Sys-*  
313 *tems*, volume 35, pages 17359–17372. Curran Associates, Inc., 2022. doi: 10.52202/  
314 068431-1262. URL [https://proceedings.neurips.cc/paper\\_files/paper/2022/  
315 hash/6f1d43d5a82a37e89b0665b33bf3a182-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2022/hash/6f1d43d5a82a37e89b0665b33bf3a182-Abstract-Conference.html).
- 316 [13] Soroush Nasiriany, Abhiram Maddukuri, Lance Zhang, Adeet Parikh, Aaron Lo, Abhishek Joshi, Ajay  
317 Mandlekar, and Yuke Zhu. RoboCasa: Large-scale simulation of household tasks for generalist robots. In  
318 *Proceedings of Robotics: Science and Systems*, Delft, Netherlands, July 2024. doi: 10.15607/RSS.2024.  
319 XX.050. URL <https://www.roboticsproceedings.org/rss20/p050.html>.
- 320 [14] NVIDIA. Cosmos 3: Omnimodal world models for physical AI. *arXiv preprint arXiv:2606.02800*,  
321 2026. doi: 10.48550/arXiv.2606.02800. URL [https://research.nvidia.com/labs/  
322 cosmos-lab/cosmos3/technical-report.pdf](https://research.nvidia.com/labs/cosmos-lab/cosmos3/technical-report.pdf).
- 323 [15] NVIDIA. Cosmos3-nano-policy-droid. Hugging Face model checkpoint and configuration, 2026. URL  
324 <https://huggingface.co/nvidia/Cosmos3-Nano-Policy-DROID>. Checkpoint revision  
325 2b9f9517.

- [16] Lu Qiu, Yizhuo Li, Yi Chen, Yuying Ge, Yixiao Ge, and Xihui Liu. Making foresight actionable: Repurposing representation alignment in world action models, 2026. URL <https://arxiv.org/abs/2606.12217>. arXiv preprint.
- [17] Sankaran Vaidyanathan, David Arbour, Aaron Mueller, Scott Niekum, and David Jensen. The curse of multiple mediators: Hidden interaction effects in activation patching, 2026. URL <https://arxiv.org/abs/2606.27510>. arXiv preprint.
- [18] Jesse Vig, Sebastian Gehrmann, Yonatan Belinkov, Sharon Qian, Daniel Nevo, Yaron Singer, and Stuart Shieber. Investigating gender bias in language models using causal mediation analysis. In *Advances in Neural Information Processing Systems*, volume 33, pages 12388–12401. Curran Associates, Inc., 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/92650b2e92217715fe312e6fa7b90d82-Abstract.html>.
- [19] Xuning Yang, Rishit Dagli, Alex Zook, Hugo Hadfield, Ankit Goyal, Stan Birchfield, Fabio Ramos, and Jonathan Tremblay. RoboLab: A high-fidelity simulation benchmark for analysis of task generalist policies. In *Proceedings of Robotics: Science and Systems*, Sydney, Australia, July 2026. doi: 10.15607/RSS.2026.XXII.096. URL <https://roboticsproceedings.org/rss22/p096.html>.
- [20] Seonghyeon Ye, Yunhao Ge, Kaiyuan Zheng, Shenyan Gao, Sihyun Yu, George Kurian, Suneel Indupuru, You Liang Tan, Chuning Zhu, Jiannan Xiang, Ayaan Malik, Kyungmin Lee, William Liang, Nadun Ranawaka, Jiasheng Gu, Yinzhen Xu, Guanzhi Wang, Fengyuan Hu, Avnish Narayan, Johan Bjorck, Jing Wang, Gwanghyun Kim, Dantong Niu, Ruijie Zheng, Yuqi Xie, Jimmy Wu, Qi Wang, Ryan Julian, Danfei Xu, Yilun Du, Yevgen Chebotar, Scott Reed, Jan Kautz, Yuke Zhu, Linxi “Jim” Fan, and Joel Jang. World action models are zero-shot policies, 2026. URL <https://arxiv.org/abs/2602.15922>. arXiv preprint.
- [21] Tianyuan Yuan, Zibin Dong, Yicheng Liu, and Hang Zhao. Fast-WAM: Do world action models need test-time future imagination?, 2026. URL <https://arxiv.org/abs/2603.16666>. arXiv preprint, version 2.
- [22] Chushan Zhang, Jinguang Tong, Xuesong Li, Yikai Wang, and Hongdong Li. Keep the future, drop the rollout: RIFT for world action models, 2026. URL <https://arxiv.org/abs/2608.11521>. arXiv preprint, version 2.
- [23] Fred Zhang and Neel Nanda. Towards best practices of activation patching in language models: Metrics and methods. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=Hf17y6u9BC>.
- [24] Weiheng Zhao, Haoyi Jiang, Xin Shi, Liu Liu, Fan Huang, Zhizhong Su, Wei Sui, and Xinggang Wang. Faster-WAM: Efficient inference-time future conditioning for robust world action models, 2026. URL <https://arxiv.org/abs/2608.04404>. arXiv preprint.

## A Reproducibility Details

The Cosmos Policy experiments use public commit 18a2acca, checkpoint revision cb689ec0, LIBERO, PyTorch 2.7 with CUDA 12.8, and two NVIDIA H100 80GB GPUs. The frozen design contains ten tasks, three state indices, eight native policy runs per state, one reciprocal donor pair per state, and three future-noise draws per direction. Confidence intervals use 10,000 bootstrap resamples of simulator states.

The Cosmos 3 experiments pin Cosmos commit 9aa98e5a, Cosmos Framework commit d4599e2e, RoboLab commit 9db0aaf0, and checkpoint revision 2b9f9517. The frozen manifest has SHA-256 46b6fe3d6abb07712d0926361587b853b2e8313a10c7faeccc0150d030e986b.

Endpoint evaluation reconstructs the environment, restores the saved state, and replays the observation prefix. Separate controls test token layout, sampler wrapping, cache record and replay, self replacement, and the action-coordinate boundary. Complete timing and aggregate compute accounting remain to be added before submission.

## B Additional Results and Boundaries

Across all 30 Cosmos Policy states, mean donor-minus-self steering is 0.186 ([0.095, 0.290]) for actions and 0.194 ([0.091, 0.305]) for endpoints. Because the three state indices may be nested within task trajectories, the first-query analysis is the cleaner confirmatory result.

The six-state RoboCasa replication is exploratory and is not pooled with LIBERO [13]. Its donor-minus-self endpoint effect is 0.0287 ([0.0101, 0.0492]), positive in 6/6 states; action steering is smaller and does not reliably exceed the Gaussian control. In the Cosmos 3 population, the two missing frozen states ended after 21 and 29 actions, before producing the 33-frame source video. Neither was replaced. The prespecified minimum of 20 states across five tasks and the factor-study minimum of ten states across five tasks were both met across six tasks.

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract, Introduction, Results, and Discussion state the future-transplantation result, its variation across decision points, and the limits on claims about planning and task success.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Result-specific qualifications appear beside the relevant findings, and the Discussion and Limitations section addresses model and state coverage, the later-state comparison, the factorized object intervention, the internal replacement experiment, and the absence of task-success results.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper reports empirical intervention results and does not claim a theorem or formal proof.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

#### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [No]

Justification: The Methods, Results, and Appendix A specify the estimands, experimental units, controls, clustering, public checkpoints, and exactness gates. The complete seed list, all repository hashes and container digests, intervention code, and run logs are not yet included.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: The source package includes machine-readable values for the paper figures, but not the intervention code, raw logs, or an anonymous public reproduction artifact.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The Methods and Results provide the state grid, branch construction, intervention targets, controls, action horizons, model interfaces, and statistical analysis needed to interpret the results. Appendix A states which exact environment details are available and which remain unreported.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]



Justification: Principal estimates have explicitly defined 95% cluster-bootstrap intervals, the independent unit and 10,000-resample procedure are stated in the Methods, and exact sign counts are reported where relevant. The text identifies controls and secondary means for which exact intervals were not available.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: Appendix A reports GPU type and software environment, but complete per-run wall-clock and aggregate compute-hour accounting is still being assembled for the submission artifact.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The experiments use public robot-policy checkpoints and simulation benchmarks, involve no people or personal data, and report negative and failed intervention results rather than selectively discarding them.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.

- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The Discussion and Limitations section discusses auditability benefits, action-redirecting misuse, false confidence from narrow simulation, and corresponding validation and safety-monitoring needs.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: This work does not release a new high-risk pretrained model or scraped dataset; it analyzes existing robot policies in simulation.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [No]

Justification: Existing assets are cited and exact source revisions are recorded, but an explicit license inventory for every upstream code and benchmark asset remains to be added before submission.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

### 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [No]

Justification: New intervention code and result manifests are documented in the working artifact, but the anonymized release package and its final license are not yet attached to this draft.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

### 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The research contains no crowdsourcing and no human participants.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.

697                   • According to the NeurIPS Code of Ethics, workers involved in data collection, cura-  
698                   tion, or other labor should be paid at least the minimum wage in the country of the  
699                   data collector.

700   **15. Institutional review board (IRB) approvals or equivalent for research with human**  
701   **subjects**

702   Question: Does the paper describe potential risks incurred by study participants, whether  
703   such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)  
704   approvals (or an equivalent approval/review based on the requirements of your country or  
705   institution) were obtained?

706   Answer: [N/A]

707   Justification: The research contains no human-subject data collection or interaction.

708   Guidelines:

709                   • The answer [N/A] means that the paper does not involve crowdsourcing nor research  
710                   with human subjects.

711                   • Depending on the country in which research is conducted, IRB approval (or equiva-  
712                   lent) may be required for any human subjects research. If you obtained IRB approval,  
713                   you should clearly state this in the paper.

714                   • We recognize that the procedures for this may vary significantly between institutions  
715                   and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the  
716                   guidelines for their institution.

717                   • For initial submissions, do not include any information that would break anonymity  
718                   (if applicable), such as the institution conducting the review.

719   **16. Declaration of LLM usage**

720   Question: Does the paper describe the usage of LLMs if it is an important, original, or  
721   non-standard component of the core methods in this research? Note that if the LLM is used  
722   only for writing, editing, or formatting purposes and does *not* impact the core methodology,  
723   scientific rigor, or originality of the research, declaration is not required.

724   Answer: [N/A]

725   Justification: Language models were used only for writing and formatting assistance, not  
726   as a component of the experimental method or scientific evaluation.

727   Guidelines:

728                   • The answer [N/A] means that the core method development in this research does not  
729                   involve LLMs as any important, original, or non-standard components.

730                   • Please refer to our LLM policy in the NeurIPS handbook for what should or should  
731                   not be described.